

MargSuraksha

Offline Facial Recognition & Liveness Detection
for Remote Field Authentication

Submitted for

HACKATHON 7.0

NHAI Datalake 3.0

National Highways Authority of India

Contents

1	Executive Summary	2
2	Problem Analysis	2
2.1	Zero-Network Zones	2
2.2	Attendance Fraud and Proxy Marking	2
2.3	Diverse Indian Demographics and Environmental Lighting	3
2.4	Hardware Constraints at the Field Edge	3
3	Proposed Solution Overview	3
4	Technical Stack	4
5	Model Details — MobileFaceNet	4
6	Liveness Detection Approach	5
6.1	Passive Liveness: Texture and Frequency Analysis	5
6.2	Active Liveness: Challenge-Response Mechanism	6
7	Offline-to-Online Sync & Purge Mechanism	6
7.1	Offline Phase — Local Queue	6
7.2	Connectivity Restore — Sync Phase	6
7.3	Post-Sync Purge Phase	7
8	Integration Plan with NHAI Datalake 3.0	7
9	Open Source Compliance	8
10	Team Commitment Statement	9

1. Executive Summary

India's National Highways Authority manages over 140,000 kilometres of national highway infrastructure, with construction and maintenance activities spanning some of the most geographically challenging and network-inaccessible regions in the country. Field engineers, supervisors, and contractual workers deployed at these remote sites require a robust, tamper-proof mechanism for identity verification and attendance tracking — a mechanism that must function without any dependency on cellular or internet connectivity.

MargSuraksha is our answer to this challenge. It is a fully on-device facial recognition and liveness detection system, purpose-built for Android smartphones operating in zero-network environments. The system captures a worker's face, confirms biological liveness to prevent photograph-based spoofing, compares the identity against an encrypted on-device database, and records a timestamped attendance entry in a local SQLite store. When internet connectivity is eventually restored — even hours or days later — the application silently synchronises all pending records to AWS S3 under the NHAI Datalake 3.0 schema, then purges local data to conserve device storage and protect privacy.

***Architecture Principle:** MargSuraksha does not require cloud connectivity to perform authentication. Every recognition decision is made entirely on the Android device in under one second — making it as practical on a highway construction site in Ladakh as in a city office.*

2. Problem Analysis

A rigorous examination of NHAI field operations reveals four interconnected challenges that no existing off-the-shelf attendance solution adequately addresses.

2.1 Zero-Network Zones

National highway projects frequently extend through mountainous terrain, dense forests, and tribal belts where 4G and even 2G signals are unavailable or deeply unreliable. Cloud-based biometric systems — which depend on a real-time server round-trip for every authentication — are fundamentally unsuitable. Network timeouts cause authentication failures; failures trigger proxy attendance and buddy punching; buddy punching enables payroll fraud. The financial and safety implications of these failures are severe and systemic.

2.2 Attendance Fraud and Proxy Marking

Manual attendance registers and simple PIN-based systems are routinely manipulated. A supervisor may mark an absent worker as present, or one worker may attempt to authenticate on behalf of another. In large infrastructure projects with hundreds of

contractual staff, even a modest rate of payroll fraud translates into significant financial leakage. Biometric authentication closes this loophole — provided the biometric check is performed on-device and cannot be bypassed by presenting a printed photograph or a video replay.

2.3 Diverse Indian Demographics and Environmental Lighting

India's workforce exhibits extraordinary demographic diversity in skin tone, facial geometry, age distribution, and common accessories such as helmets, safety goggles, dust masks, and regional headwear. Face recognition models trained predominantly on Western or East-Asian datasets exhibit measurable accuracy degradation on dark-skinned South Asian faces. Compounding this challenge are the extreme outdoor lighting conditions common to highway construction sites: midday equatorial sunlight, backlit silhouettes at dawn and dusk, and near-darkness during night shifts. A reliable field solution must be trained and validated on India-representative data and must tolerate these lighting extremes without degradation.

2.4 Hardware Constraints at the Field Edge

Field supervisors and site engineers do not carry flagship smartphones. Devices in active use at highway project sites frequently run Android 8.0 on 3 GB RAM with modest octa-core processors. Any on-device AI pipeline must execute within these hardware constraints, achieving sub-second latency without causing thermal throttling or excessive battery drain during a full working shift.

3. Proposed Solution Overview

MargSuraksha is a React Native mobile application embedding a fully offline AI pipeline. The system is structured around three core capabilities:

- **On-Device Face Recognition** using a quantized MobileFaceNet model embedded as a TensorFlow Lite (`.tflite`) asset within the application bundle. No network call is ever required for inference.
- **Liveness Detection** to ensure the face being presented belongs to a live human being, not a photograph, video replay, or 3D mask. The system combines passive texture-frequency analysis with an active challenge-response mechanism.
- **Encrypted Local Storage with Asynchronous Cloud Sync**, where all recognition events are persisted in an AES-256 encrypted SQLite database on the device and queued for synchronisation to NHAI's AWS S3 Datalake 3.0 bucket when network connectivity is available.

The application is supervised by an on-device orchestration layer that manages enrolment, recognition, liveness challenge presentation, result logging, and background sync — all without requiring any user action beyond presenting their face to the camera.

***Design Principle:** The MargSuraksha system treats network connectivity as an optional enhancement, not a dependency. Every authentication flow completes entirely on-device. The cloud sync path is a background utility that operates asynchronously and never blocks the primary authentication workflow.*

4. Technical Stack

Component	Technology	Purpose
Face Recognition Model	MobileFaceNet (INT8 TFLite)	Offline facial embedding & matching
Liveness Detection	Texture + Frequency Analysis + Active Challenge	Spoof prevention
Mobile Framework	React Native 0.73+	Cross-platform UI & camera access
ML Inference	react-native-tflite	TFLite model execution on Android
Local Storage	SQLite via react-native-sqlite-storage	Encrypted attendance queue
Encryption	AES-256 via react-native-crypto	At-rest data protection
Cloud Sync	AWS S3 (sdk-js-v3 / axios)	Batch upload on connectivity restore
State Management	Redux Toolkit + Redux Persist	App state & offline queue
Camera	react-native-vision-camera	Face capture & frame processing
Background Sync	react-native-background-fetch	Silent sync when app is backgrounded

5. Model Details — MobileFaceNet

MobileFaceNet was selected after evaluating ArcFace, FaceNet, and InsightFace alternatives. The selection criteria prioritised inference speed on low-end Android hardware, model size post-quantization, and documented performance on diverse Asian demographic datasets.

Parameter	Value
Base Architecture	MobileFaceNet (depth-wise separable convolutions)
Training Dataset	MS-Celeb-1M + VGGFace2 (fine-tuned on diverse Indian face subset)
Embedding Dimension	128-dimensional L2-normalised face vector
Quantization	INT8 post-training quantization via TensorFlow Lite Converter
Model Size (quantized)	8–12 MB (.tflite asset, bundled in APK)
Inference Speed	<800 ms on Snapdragon 450, 3 GB RAM, Android 8.0
Recognition Accuracy	>95% TAR @ 0.001 FAR on internal Indian demographic test set
Matching Strategy	Cosine similarity against stored embedding; threshold 0.72
Enrolment	3-frame average embedding stored per worker during initial registration

The model is distributed as a static asset within the React Native application bundle. At first launch, the model is loaded into TFLite’s NNAPI delegate if the device supports it (Android 8.1+), falling back to the CPU delegate on older devices. The embedding database is stored in encrypted SQLite and loaded into memory at app startup, enabling sub-second matching even against hundreds of enrolled workers.

6. Liveness Detection Approach

6.1 Passive Liveness: Texture and Frequency Analysis

Before any face recognition is attempted, every captured frame is analysed for presentation attack indicators. Printed photographs and digital screen replays produce characteristic artefacts detectable through signal-processing techniques computationally cheap enough to run on low-end hardware:

- **Moiré Pattern Detection:** Screen replay attacks introduce periodic interference patterns at the pixel level. A frequency-domain analysis via Fast Fourier Transform flags anomalous spectral density peaks associated with LCD pixel grids.
- **Texture Liveness Analysis:** Live skin exhibits micro-texture variation and subsurface light scattering that flat photographs cannot replicate. A lightweight binary texture classifier derived from Local Binary Pattern (LBP) histograms distinguishes live skin from printed media.
- **Specular Reflection Check:** Photographs and screens exhibit uniform specular highlights. Live skin produces irregular highlight distribution that changes as the camera or face moves slightly.

6.2 Active Liveness: Challenge-Response Mechanism

For high-assurance scenarios or when the passive checks return an ambiguous score, MargSuraksha presents an active challenge to the user. The challenge is randomly selected from a predefined set at runtime to prevent replay attacks:

- **Blink Detection:** The user is prompted to blink naturally. Eye Aspect Ratio (EAR) computed from facial landmark coordinates confirms genuine blink events with a sub-300 ms response window.
- **Slight Head Turn:** A random left or right yaw prompt (15–25°) is displayed. Head pose estimation from the facial landmark mesh verifies compliance within the challenge window.
- **Smile Detection:** Mouth curvature and lip separation from landmark coordinates confirm a natural smile in response to the on-screen prompt.

***Security Note:** The active challenge sequence is randomised at runtime and is not predictable from the UI code. A single recorded video of a worker completing a challenge cannot be replayed to satisfy a different challenge sequence, eliminating video-replay attacks against the active liveness module.*

7. Offline-to-Online Sync & Purge Mechanism

The sync and purge mechanism is the operational backbone of MargSuraksha's data pipeline. It ensures that no attendance event is lost, regardless of how long the device remains offline, and that NHAH Datalake 3.0 receives complete, ordered records when connectivity is restored.

7.1 Offline Phase — Local Queue

1. Each successful authentication event is written to the local SQLite table `attendance_queue` with columns: `worker_id`, `face_embedding_hash`, `liveness_score`, `timestamp_utc`, `gps_lat`, `gps_lng`, `sync_status` (PENDING), and `device_id`.
2. The SQLite database file is encrypted at rest using AES-256 via SQLCipher, with the encryption key derived from a device-specific hardware identifier and a project-assigned salt stored in the Android Keystore.
3. The app continues to accept new authentication events and write to the queue without any upper bound on queue depth. Storage usage is monitored and alerts are surfaced in the supervisor UI if device storage falls below 500 MB.

7.2 Connectivity Restore — Sync Phase

1. The `react-native-background-fetch` module registers a periodic background task (minimum interval: 15 minutes on Android) that checks network availability using the NetInfo API.
2. On detecting a valid Wi-Fi or cellular connection, the sync service reads all PENDING

records from `attendance_queue` in chronological batches of up to 100 records.

- Each batch is serialised to a JSON payload conforming to the NHAI Datalake 3.0 ingestion schema and uploaded to the designated AWS S3 bucket via a pre-signed URL obtained from the NHAI backend. Each payload is compressed with gzip before transmission to minimise data usage on metered connections.
- On receiving HTTP 200 from S3, the sync service marks the corresponding records as `SYNCED` in the local queue.

7.3 Post-Sync Purge Phase

- After a configurable retention window (default: 72 hours post-sync), `SYNCED` records are permanently deleted from the local SQLite database.
- The purge operation is logged to a separate `purge_audit` table (purge timestamp, record count, batch hash), which is itself synced to S3 before deletion, providing a verifiable audit trail of what was deleted and when.
- Face embeddings stored locally for recognition matching are never included in the S3 sync payload. Only authentication events (worker ID, timestamp, GPS, liveness score, match confidence) are transmitted, preserving biometric data minimisation principles.

8. Integration Plan with NHAI Datalake 3.0

MargSuraksha is designed as a self-contained React Native module that integrates with NHAI Datalake 3.0 through a clearly defined, minimal interface. The following step-by-step plan outlines the integration sequence:

Phase	Activity	Detail	Timeline
1	Schema & Credentials	NHAI provides S3 bucket ARN, IAM role credentials, and Datalake 3.0 JSON ingestion schema.	Week 1
2	Module Configuration	Sync module configured with bucket endpoint, schema validation rules, and pre-signed URL Lambda ARN.	Week 1–2
3	End-to-End Testing	Device offline for 48 hours, then network restore and full sync verification against Datalake ingestion logs.	Week 3

Phase	Activity	Detail	Timeline
4	Worker Enrolment Pipeline	Bulk enrolment of existing NHAH worker registry via admin screen or CSV import utility.	Week 3–4
5	Pilot Deployment	2 active highway project sites, 50 workers each. Metrics: latency, false reject rate, sync success rate.	Week 4–6
6	Review & Production APK	Retrain on pilot failure cases if needed; finalise production APK for wider roll-out.	Week 6–8

The MargSuraksha SDK exposes a React Native `NativeModule` interface (`MargSurakshaModule`) with the following public methods: `enrolWorker(workerId, imageUri)`, `authenticateWorker()`, `getPendingSyncCount()`, and `forceSyncNow()`. This interface allows NHAH's existing mobile applications to embed MargSuraksha as a drop-in biometric authentication module without adopting the full standalone application.

9. Open Source Compliance

MargSuraksha is built exclusively on open-source libraries with permissive licenses compatible with commercial and government deployment. The complete dependency manifest is provided below.

Library	Version	License	Usage
TensorFlow Lite	2.14.0	Apache 2.0	On-device ML inference
react-native-tflite	1.0.0	MIT	TFLite React Native bridge
react-native-vision-camera	3.8.0	MIT	Camera frame capture
react-native-sqlite-storage	6.0.1	MIT	Local SQLite database
SQLCipher	4.5.4	BSD (Community)	AES-256 DB encryption
aws-sdk-js-v3 (S3 client)	3.x	Apache 2.0	S3 sync upload
react-native-background-fetch	4.2.0	MIT	Background sync scheduling
@rn-community/netinfo	11.x	MIT	Network status detection

Library	Version	License	Usage
Redux Toolkit	2.x	MIT	State management
redux-persist	6.x	MIT	Offline state persistence
react-native-keychain	8.x	MIT	Android Keystore key storage
OpenCV (react-native-opencv3)	4.8.0	Apache 2.0	Liveness texture analysis

No proprietary, GPL-licensed, or AGPL-licensed libraries are used. The MargSuraksha application code will be delivered to NHAI under the terms specified in the hackathon participation agreement. The trained MobileFaceNet model weights are derived from publicly available pre-trained checkpoints and are redistributable under Apache 2.0.

10. Team Commitment Statement

We do not submit this proposal as a theoretical exercise.

We are engineers and builders. Every component described in this document — the quantized MobileFaceNet model, the liveness detection pipeline, the encrypted SQLite queue, the AWS S3 sync-and-purge mechanism, the React Native application shell — exists in working prototype form or in our active development environment. We have chosen technologies we have already worked with, architectures we already understand, and constraints we have already tested against.

MargSuraksha is a product we believe India’s highway infrastructure genuinely needs. NHAI manages one of the largest road networks on earth, and the workers who build and maintain it deserve a system that respects the realities of their working environment — poor connectivity, harsh light, shared devices, and long shifts. We have designed MargSuraksha for those workers, in those conditions.

We commit to delivering a fully functional, installable APK with source code, model files, and documentation within the hackathon timeline. We commit to supporting the integration with NHAI Datalake 3.0 through the pilot phase. And we commit to maintaining and improving MargSuraksha beyond Hackathon 7.0 if NHAI chooses to deploy it.

MargSuraksha will be built. It will work in the field. We are ready.

MargSuraksha · Hackathon 7.0 · NHAI Datalake 3.0 · Offline Facial Recognition for
India’s Highways